

Iteration-5 confirmation arms: status, AUROC and verdict

Arm (labelled set)	Status	Consensus V0 AUROC [95% CI]	Judge (disguised) AUROC [95% CI]	Pre-registered verdict	Arm spend
E2-A 110 / 550 sentences panel-labelled	DT: testable 179 rows 74 ERR / 105 COR	0.916 [0.858, 0.970] n = 179	0.913 [0.858, 0.958] Qwen3-8B · n = 179	NOT_TESTABLE flash-lite judge bar missing	\$2.05
	L25: not testable 42 rows 35 ERR / 7 COR	0.443 [0.375, 0.500] n = 42	0.659 [0.450, 0.967] Qwen3-8B · n = 42		
E2-B 400 sentences no panel labels	48 CORRECT < 50 334 ERR / 48 COR (vs unaudited gold)	0.737 [0.645, 0.827] exploratory, n = 381	0.496 [0.356, 0.620] flash-lite · n = 276	NOT_TESTABLE CORRECT floor not met	\$0.83
R_COMP FREE 2,265 labelled rows	Gloss gate FAILED ×2 759 ERR / 1,506 MAPPED → 0 CORRECT	0.801 [0.770, 0.834] provisional, n = 1,706	0.584 [0.550, 0.618] flash-lite · n = 1,706	NOT_TESTABLE 0 CORRECT rows (fallback B)	\$1.97

Shared \$12.00 OpenRouter budget
exhausted (~11:40 UTC)

Testable cell
 Exploratory / provisional (cannot confirm)
 Not testable